

VIHARI REDDY ALETI

aleti.viharireddy@gmail.com | Dallas, TX (open to relocation) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Master's graduate in Information Technology with experience in data driven consulting. Built metadata driven ETL pipelines, and CDC ingestion frameworks using Python, SQL, Spark, and cloud technologies. Developed scalable data solutions that improve data reliability, enable analytics and business decision making.

EDUCATION

Master of Science in Information Technology

Belhaven University

Jan 2025 - Apr 2026

Mississippi

Bachelor of Engineering in Computer Science

Jntuh

Aug 2016 - Sep 2020

Hyderabad, India

SKILLS

- **Languages:** Python, SQL, Java, Bash
- **Big Data & Libraries:** Apache Spark, PySpark, Hive, Pandas, NumPy
- **Data Engineering:** Apache Airflow, dbt, ELT/ETL Pipelines, Data Cleaning
- **Databases & Warehousing:** PostgreSQL, MySQL, Amazon Redshift, Snowflake
- **Cloud & DevOps:** AWS (S3, EC2), Azure, Docker, Kubernetes
- **Machine Learning & Gen AI:** Scikit-learn, MLflow, NLP, LLM, LangChain, RAG

PROJECTS

Agentic Aegis AI Platform | [Link](#)

Tools Used: Python, SQL, Polars, DuckDB, AWS, Prefect, Docker, Scikit-learn, LLM

- Architected a self healing ETL platform with 4 autonomous agents for data profiling, remediation, and anomaly detection across 8 validation categories, cutting manual intervention 60%
- Built a Bronze-Silver-Gold data warehouse using DuckDB, Polars, Prefect with automated pipelines, enabling full lineage auditability and reliable analytics
- Developed an LLM-powered natural language to SQL interface with schema aware prompting and query validation, restricting execution to read only queries
- Automated containerized deployment, orchestration, and monitoring using Docker, Kubernetes, and Prometheus, reducing environment setup time and deployment effort

Lakeflow Data Ingestion Framework | [Link](#)

Tools Used: Python, PySpark, Apache Spark, Databricks, Delta Lake, PostgreSQL, CI/CD

- Built a metadata driven ingestion framework enabling full-load and append only CDC sync from PostgreSQL to Databricks Delta Lake using WAL, dlt, and Databricks Lakeflow.
- Added automated data quality checks, retry logic, and structured alerting, cutting ingestion failure impact and improving pipeline reliability
- Automated CI/CD with Databricks Asset Bundles and GitHub Actions, eliminating manual releases via end-to-end build, test, deploy workflows
- Built integration tests and reusable Python ingestion components for full load and CDC pipelines, cutting manual ETL effort by 80% and accelerating production readiness
- Validated PostgreSQL WAL-based CDC end-to-end, verifying INSERT, UPDATE, and DELETE propagation and soft-delete correctness in Delta Lake.

WORK EXPERIENCE

Executive Consultant | IT World Web | Dec 2022 - Dec 2024 | Bangalore, India

- Analyzed 10K+ operational records across 5 business functions to identify process bottlenecks, map current-state workflows, and recommend changes that improved operational efficiency by 40%.
- Established and standardized 10+ enterprise KPIs by translating business requirements into measurable performance indicators, improving reporting accuracy and strategic planning.

Associate Consultant | Careernet Technologies | Apr 2021 - Nov 2022 | Bangalore, India

- Built executive dashboards by consolidating operational metrics from multiple functions, enabling faster performance tracking and improved leadership visibility.
- Investigated data and process level discrepancies through root cause analysis, identifying performance drivers and delivering actionable insights to cross-functional stakeholders.

CERTIFICATIONS

- Data Engineering on AWS Foundations | **AWS**
- Analytical SQL for Developers | **Oracle**